

Computer Architectures

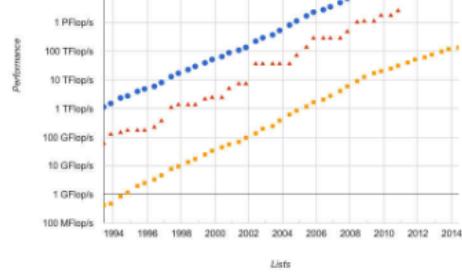
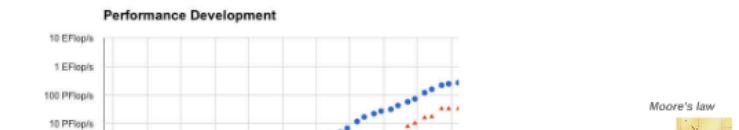
Scientific Computing
Jörn Behrens



Introduction

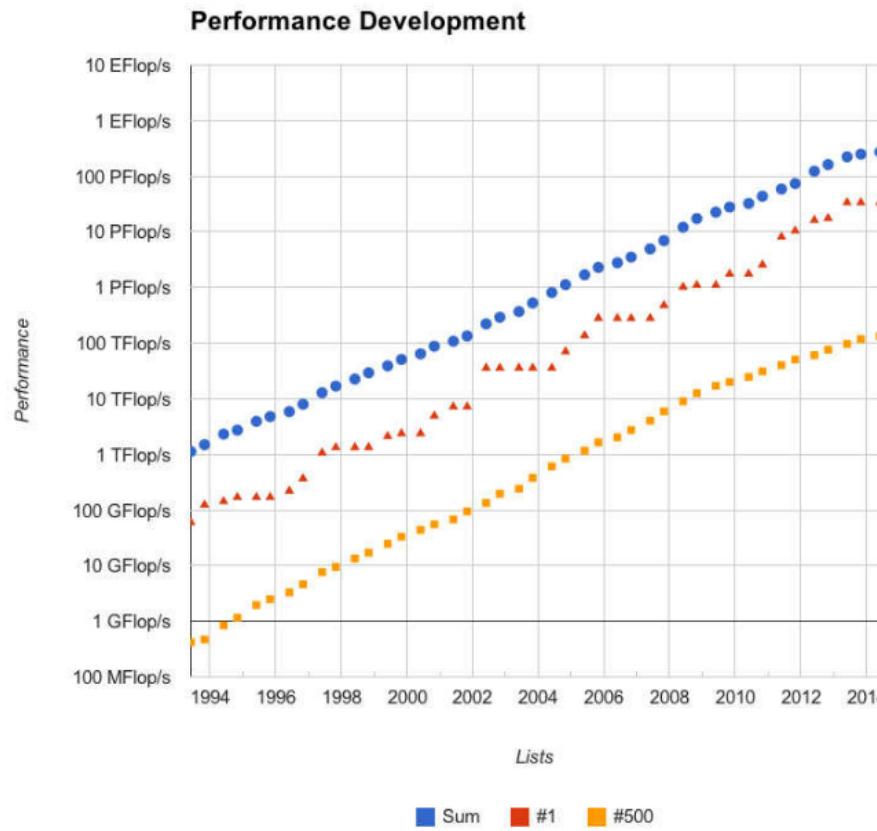
Why deal with Computer Architectures?

- Most programs achieve 10-20% of peak performance, why?
- How does a computer work, anyways?
- What makes the performance?

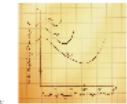


Why deal with Computer Architectures?

- Most programs achieve 10-20% of peak performance, why?
- How does a computer work, anyways?
- What makes the performance?



Moore's law



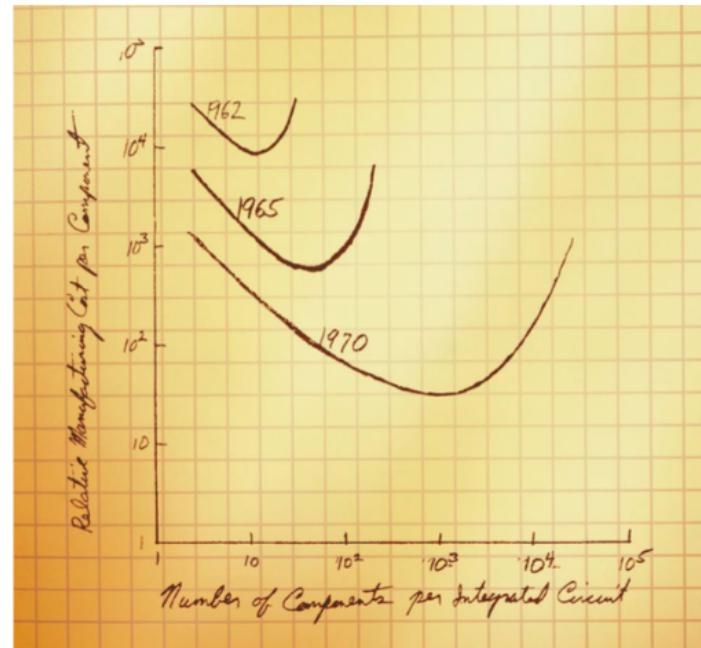
Original:
every 18 months:
double components per integrated circuit

every 18 months:
double computational performance!

Moore's law

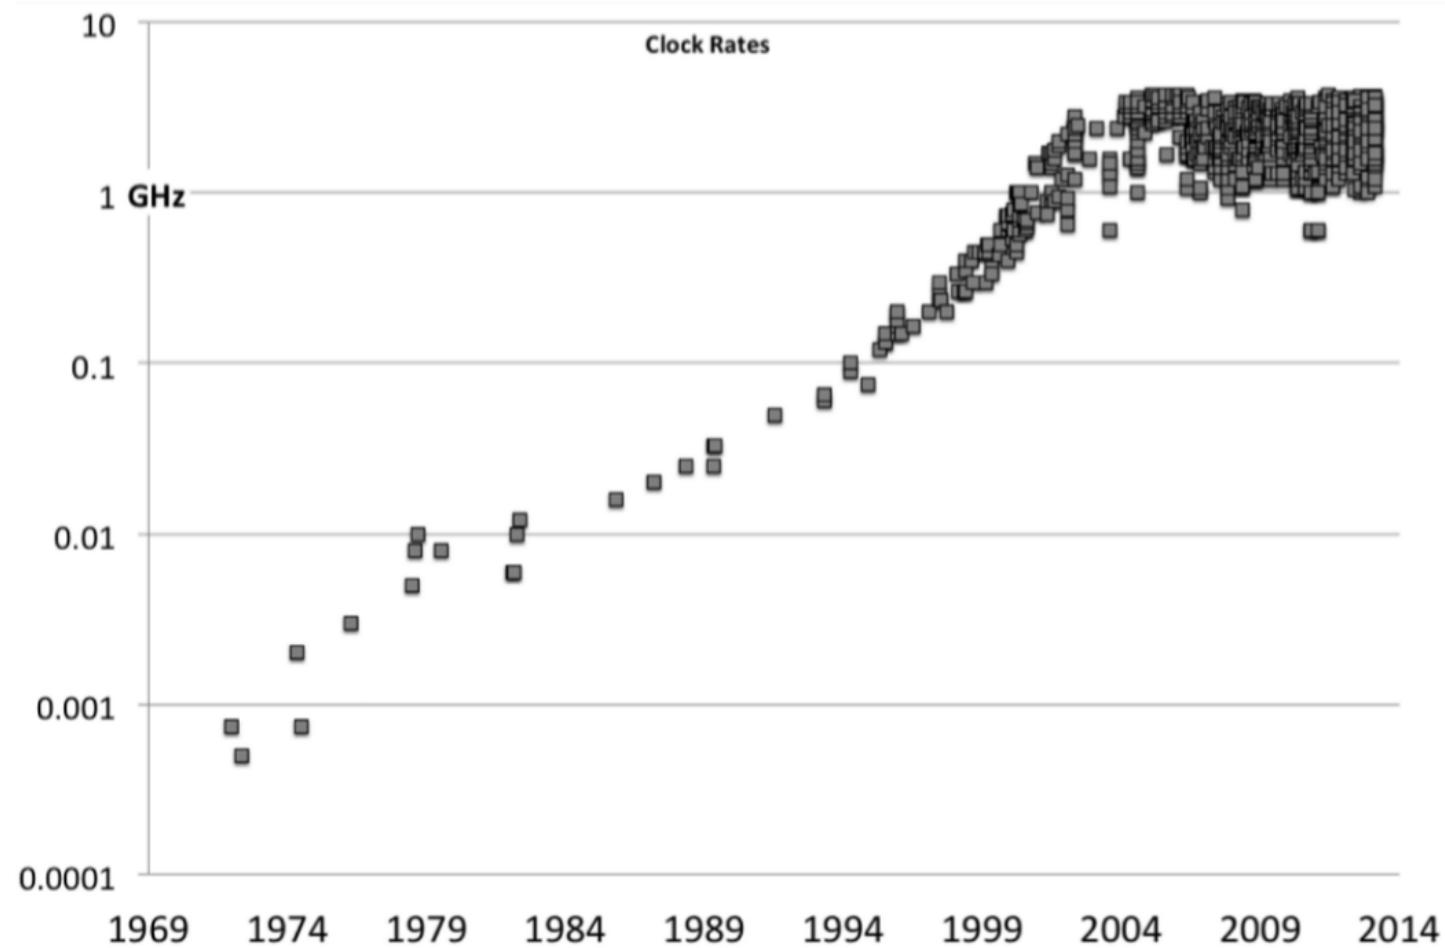
Original:

every 18 months:
double components per integrated circuit

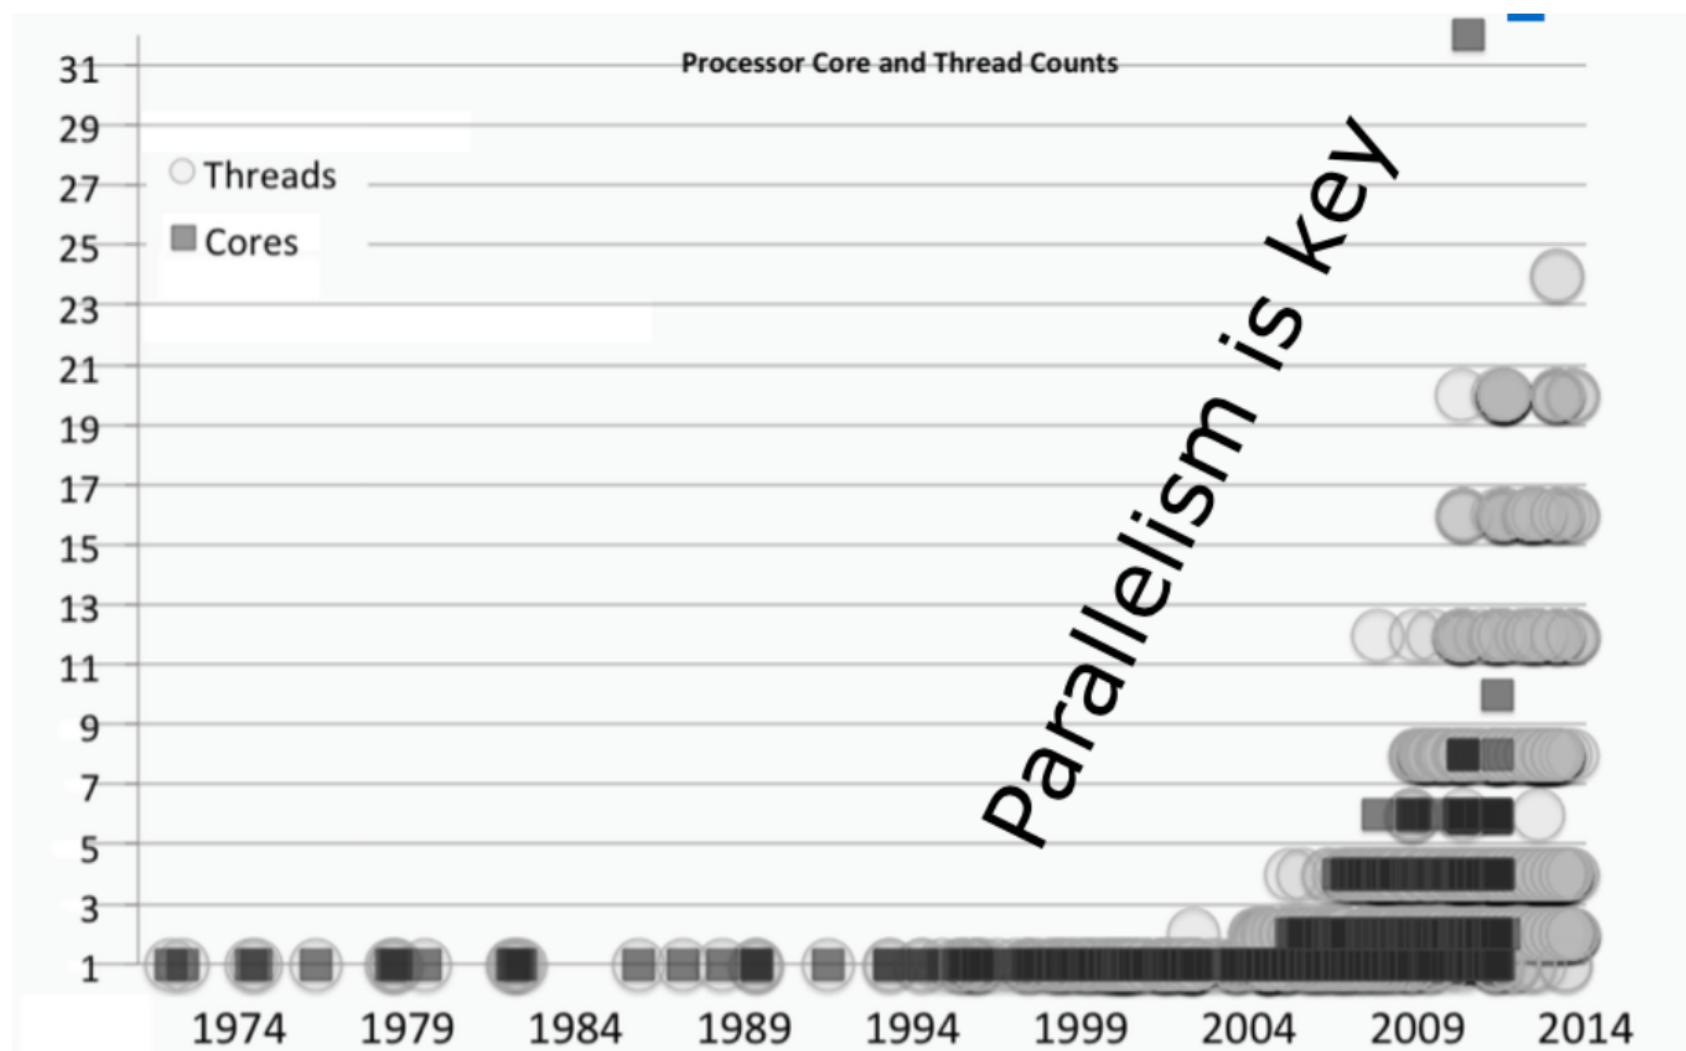


every 18 months:
double computational performance!

Increase in Clockrate?

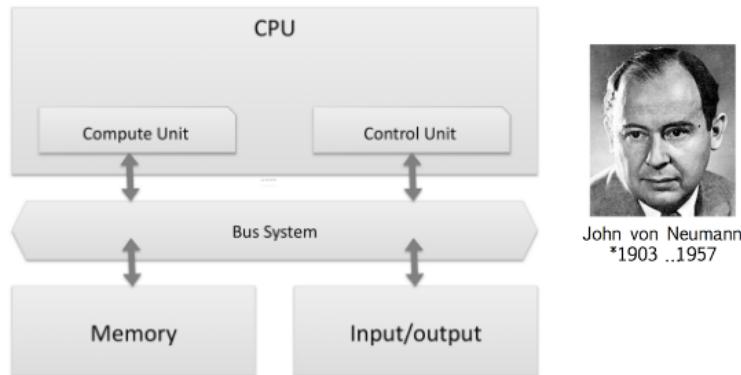


Parallelism is the Key



Basics

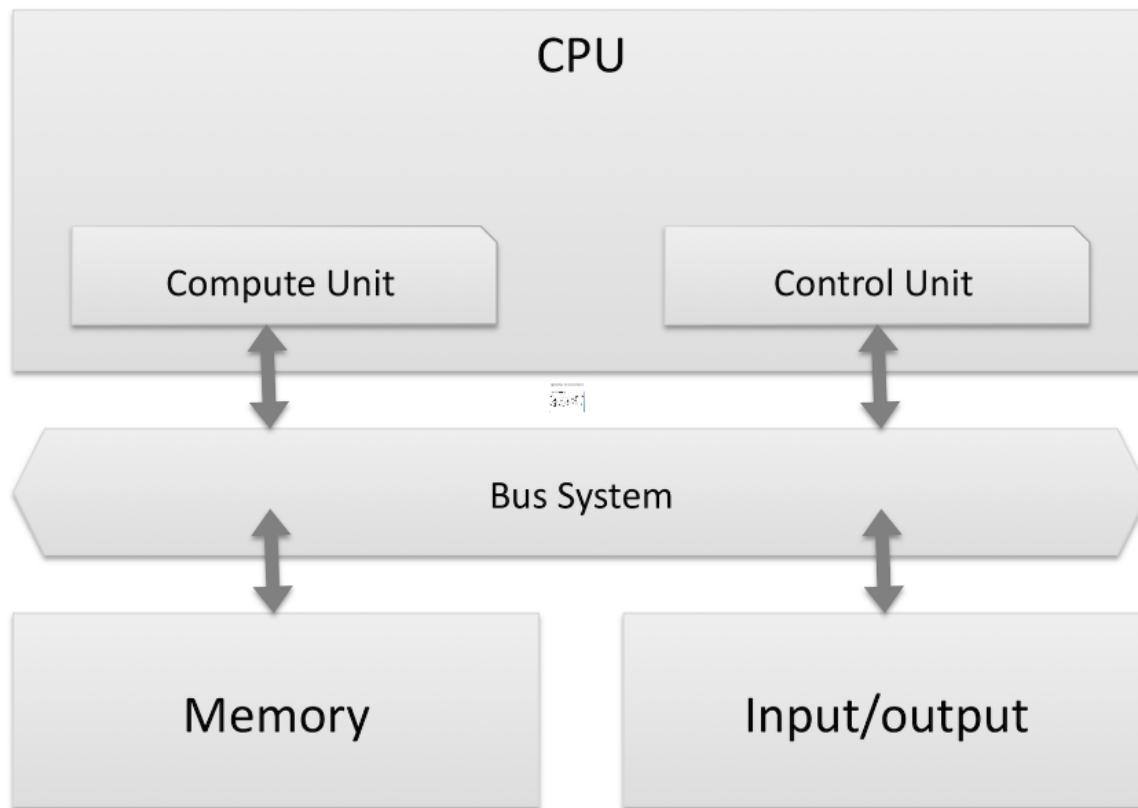
Von Neumann Architecture



John von Neumann
*1903 ..1957

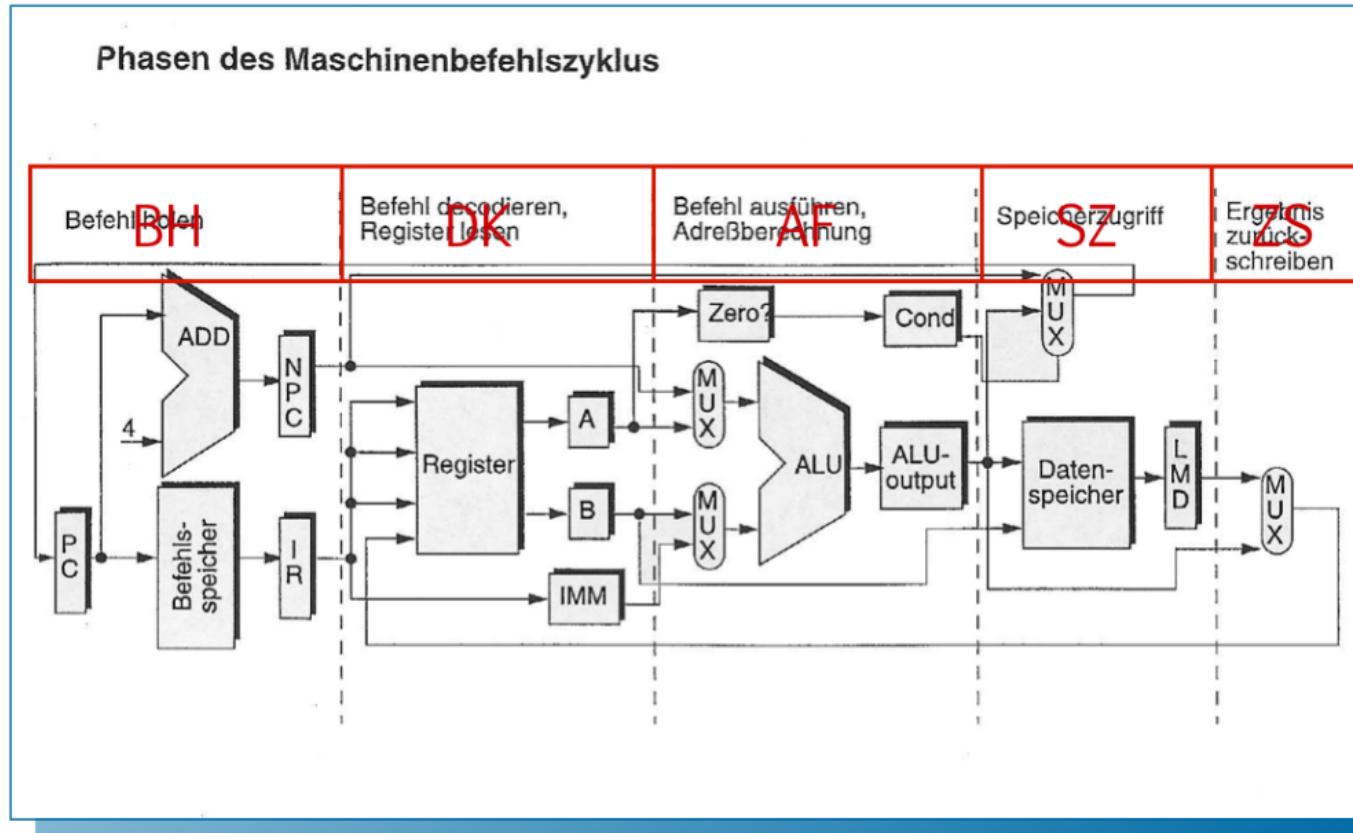


Von Neumann Architecture

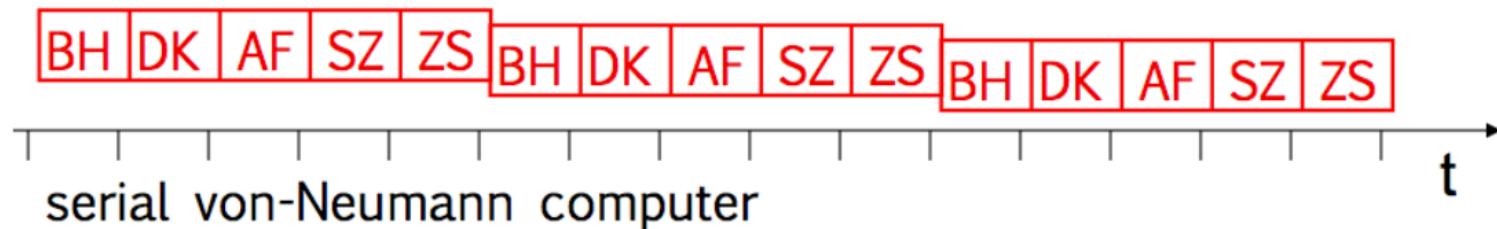


John von Neumann
*1903 ..1957

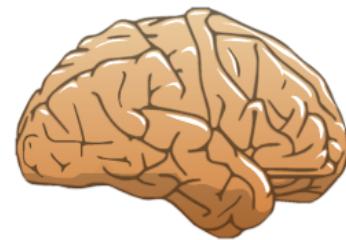
Operation via Instructions



How to accelerate?

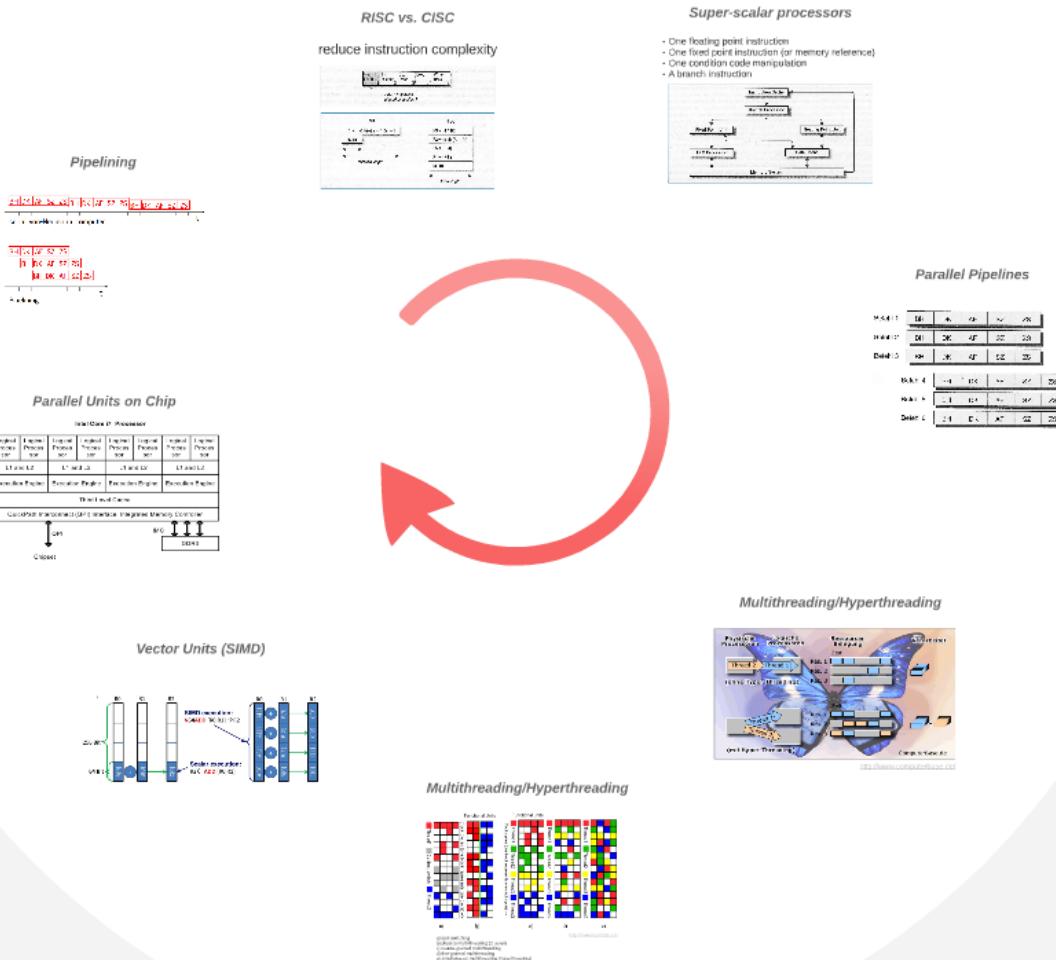


What are your ideas?

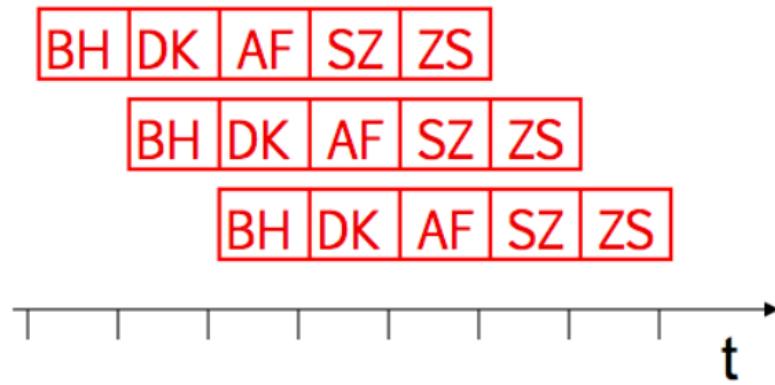
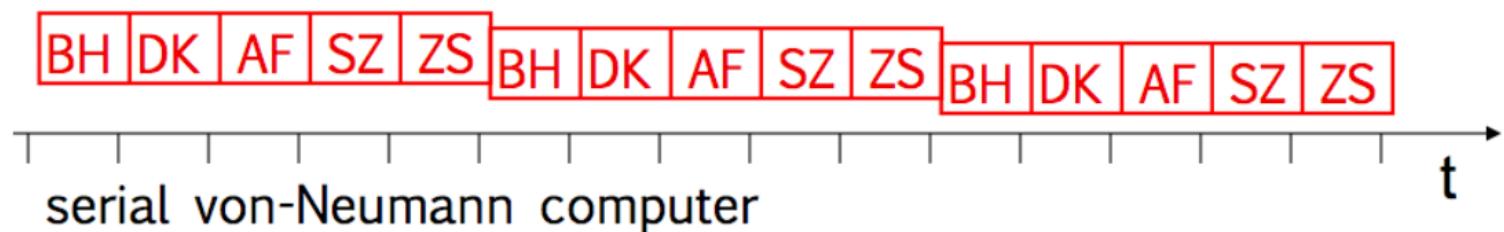


Acceleration Features

Parallelism on the processor level

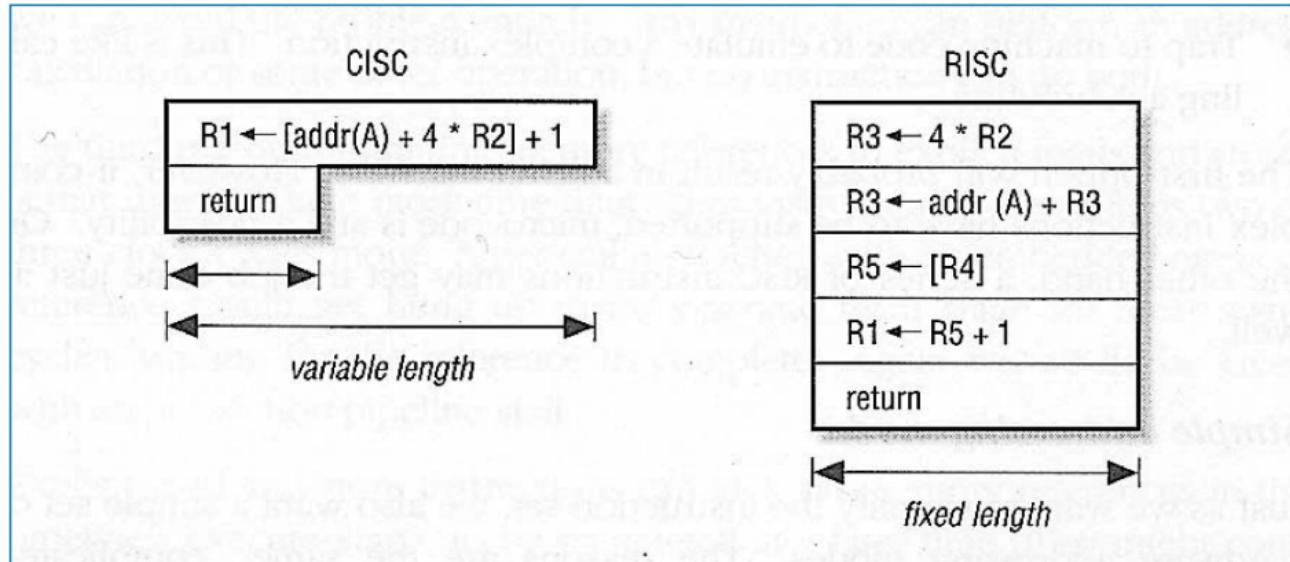
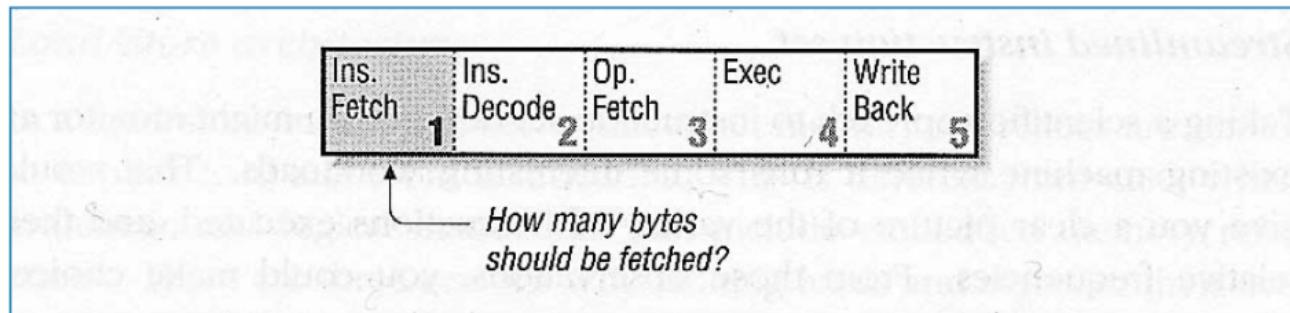


Pipelining



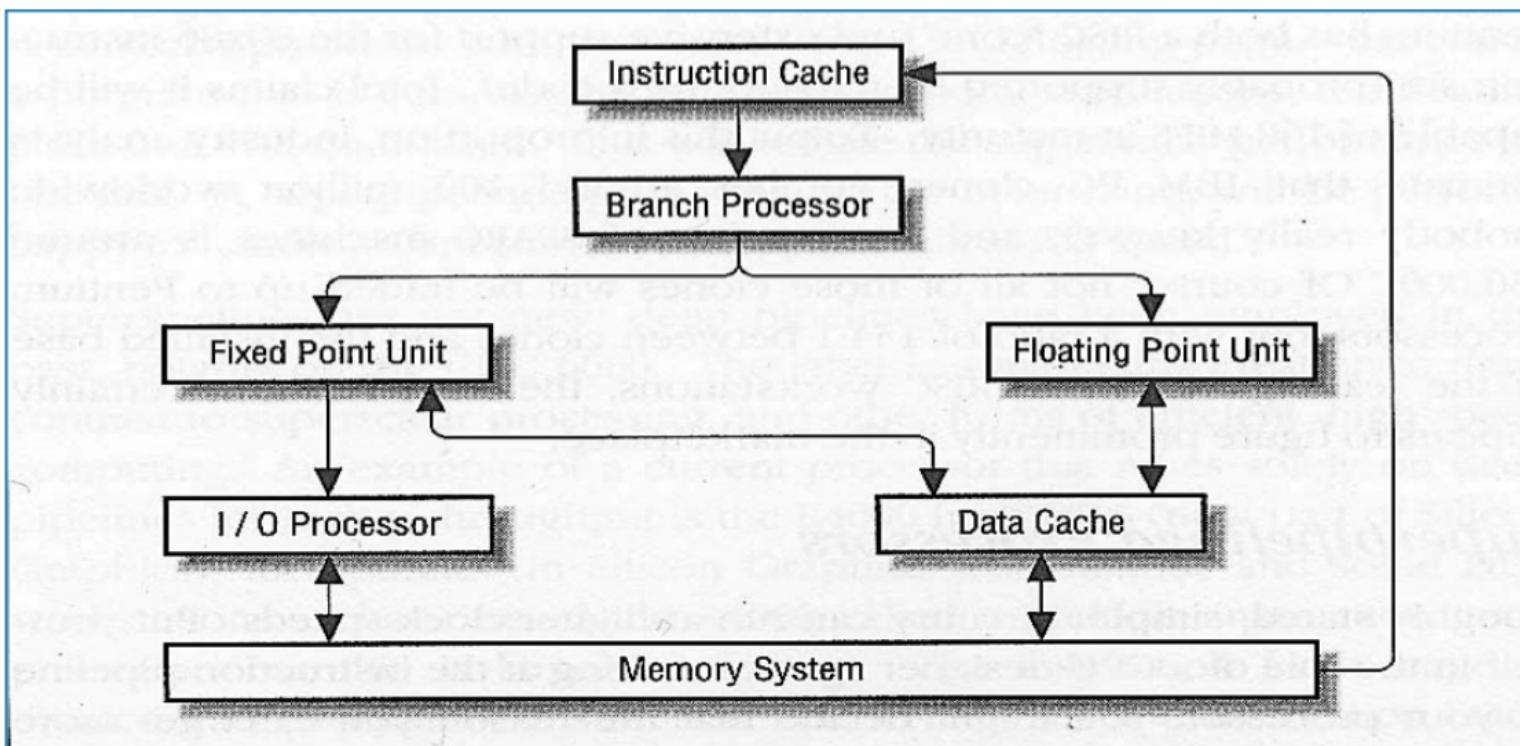
RISC vs. CISC

reduce instruction complexity

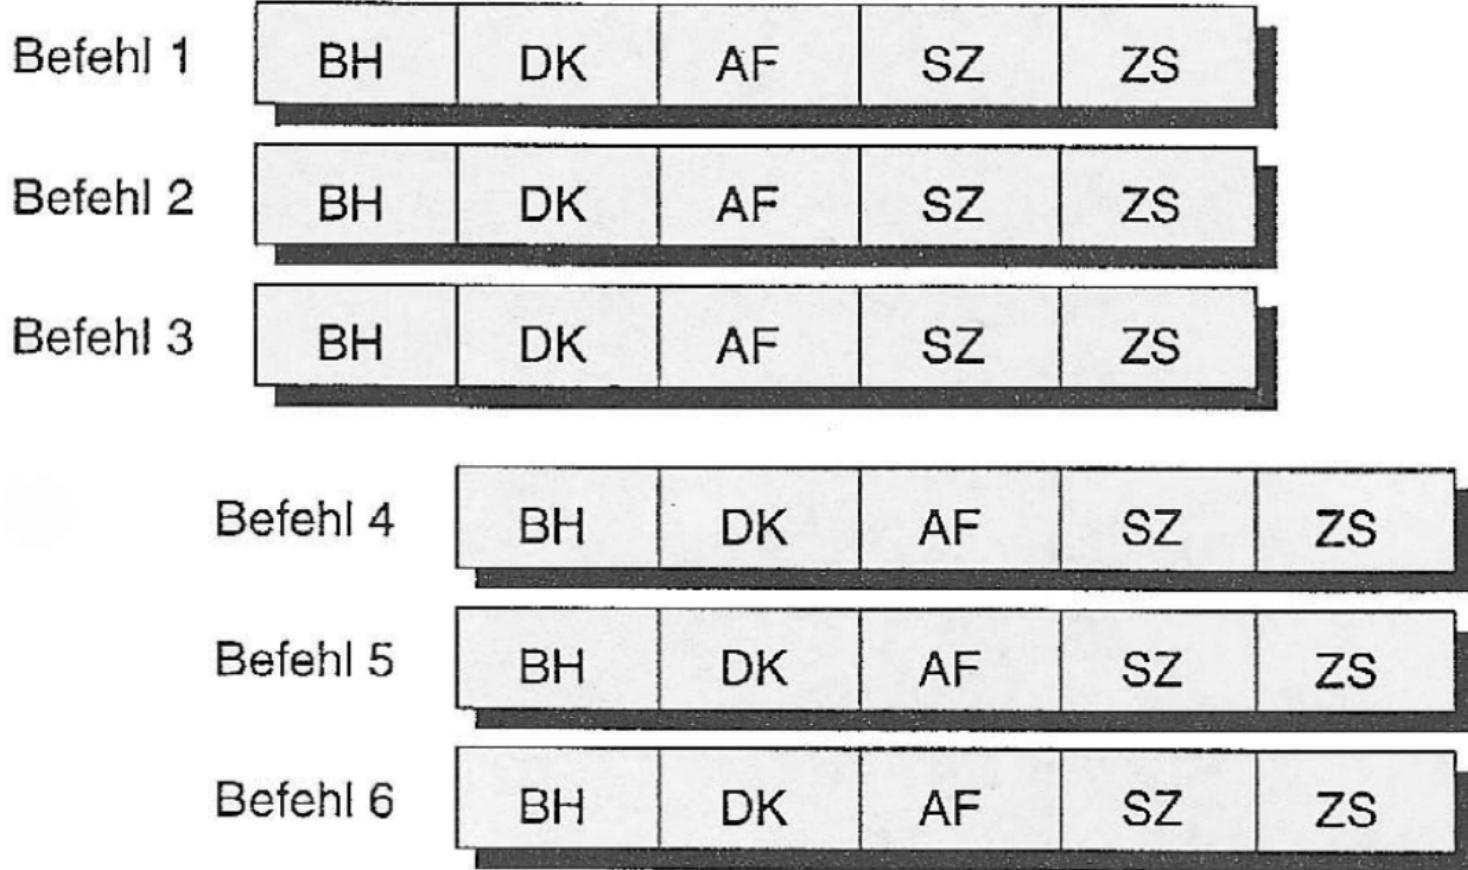


Super-scalar processors

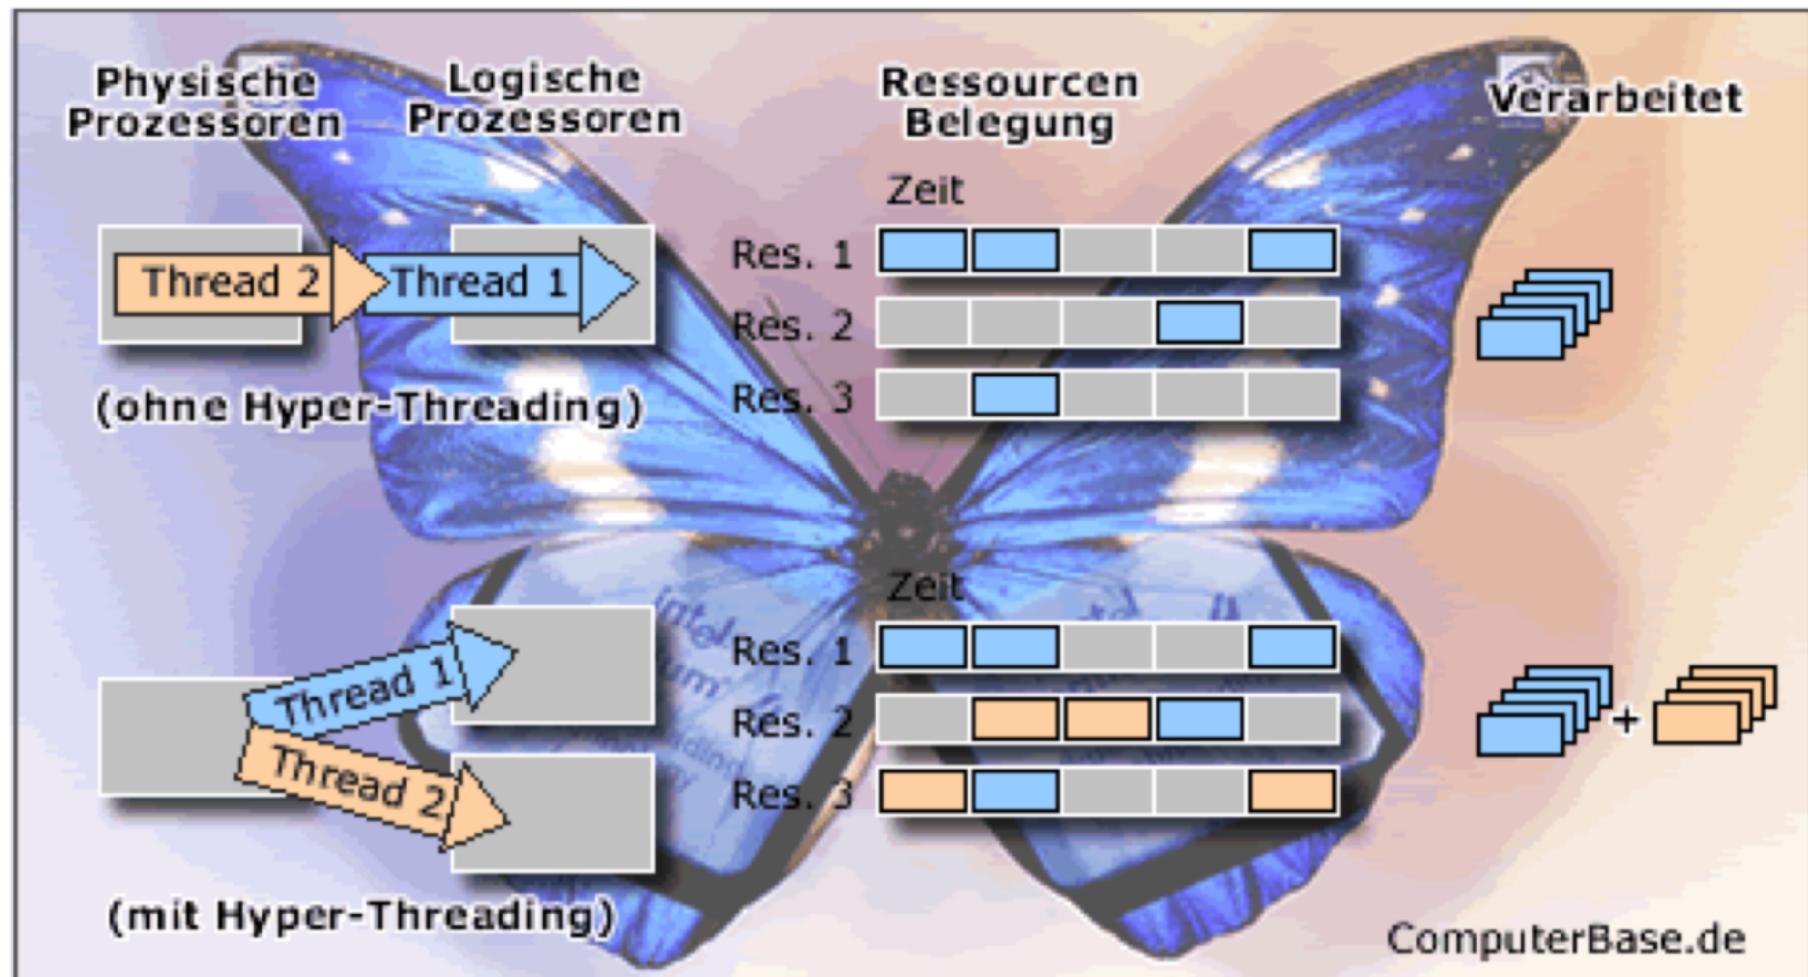
- One floating point instruction
- One fixed point instruction (or memory reference)
- One condition code manipulation
- A branch instruction



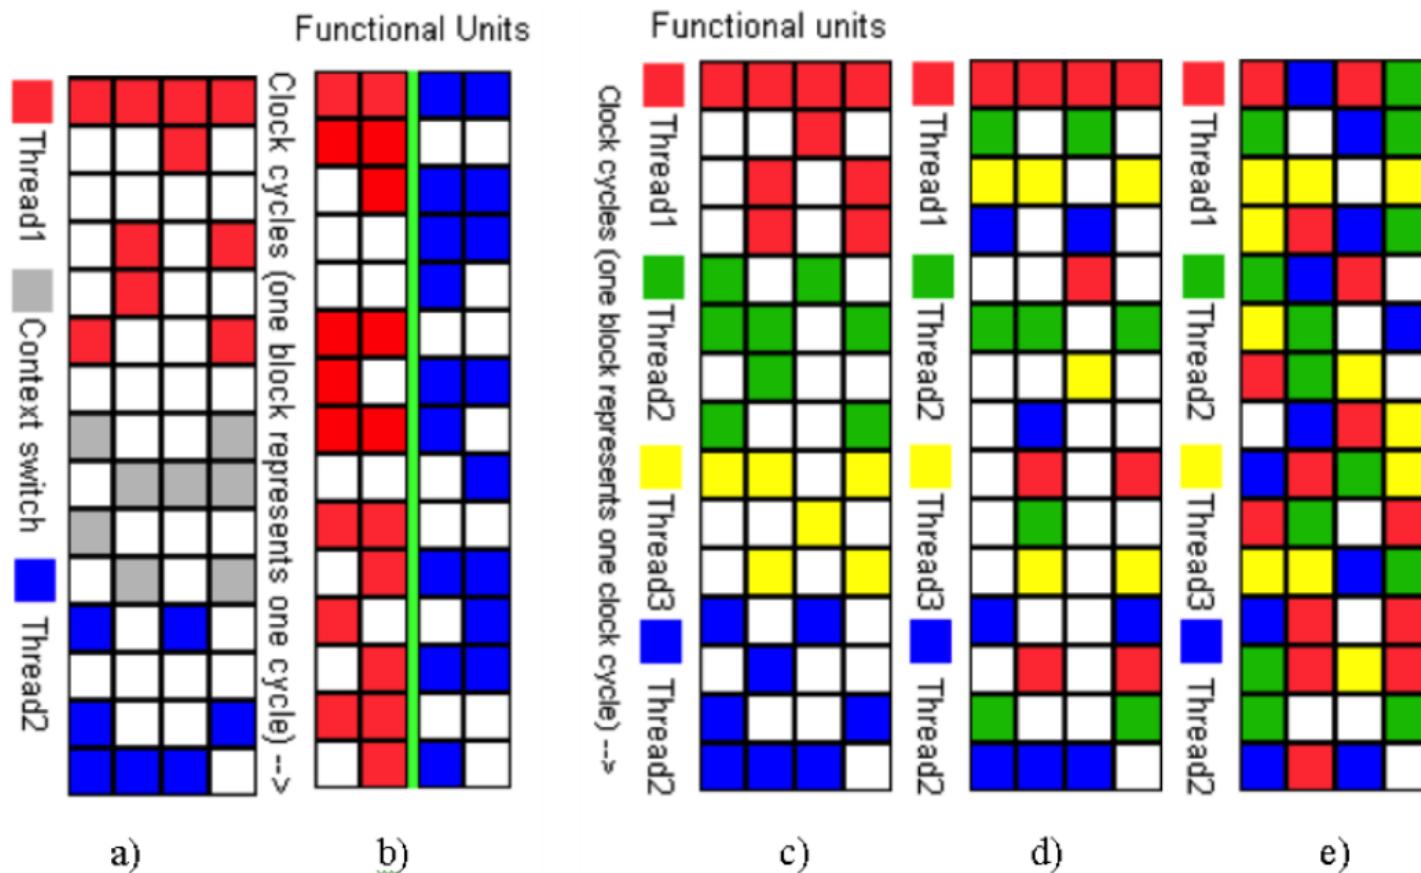
Parallel Pipelines



Multithreading/Hyperthreading



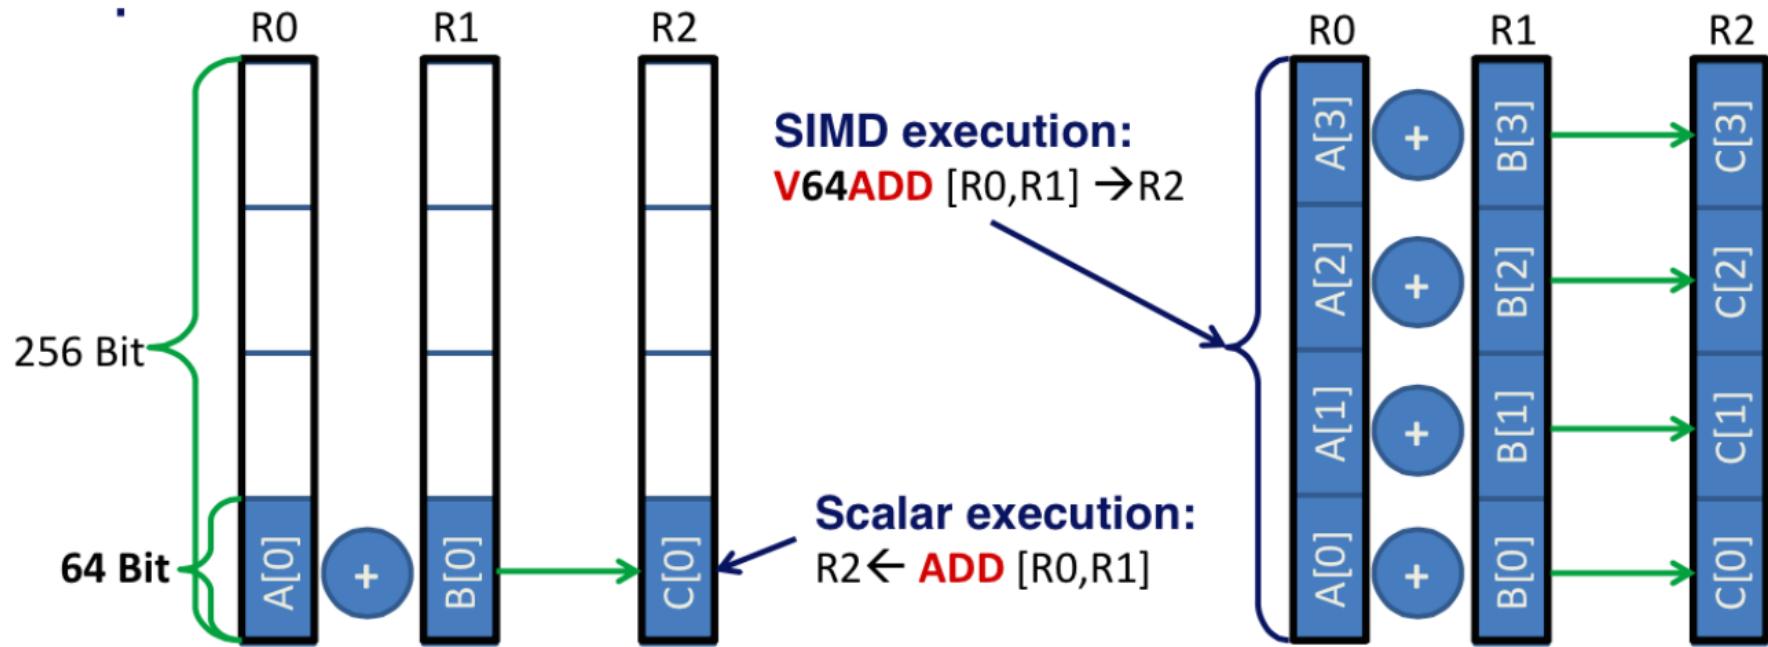
Multithreading/Hyperthreading



- a) task switching
- b) physical multi-threading (2 cores)
- c) coarse-grained multithreading
- d) fine-grained multithreading
- e) simultaneous multithreading (hyperthreading)

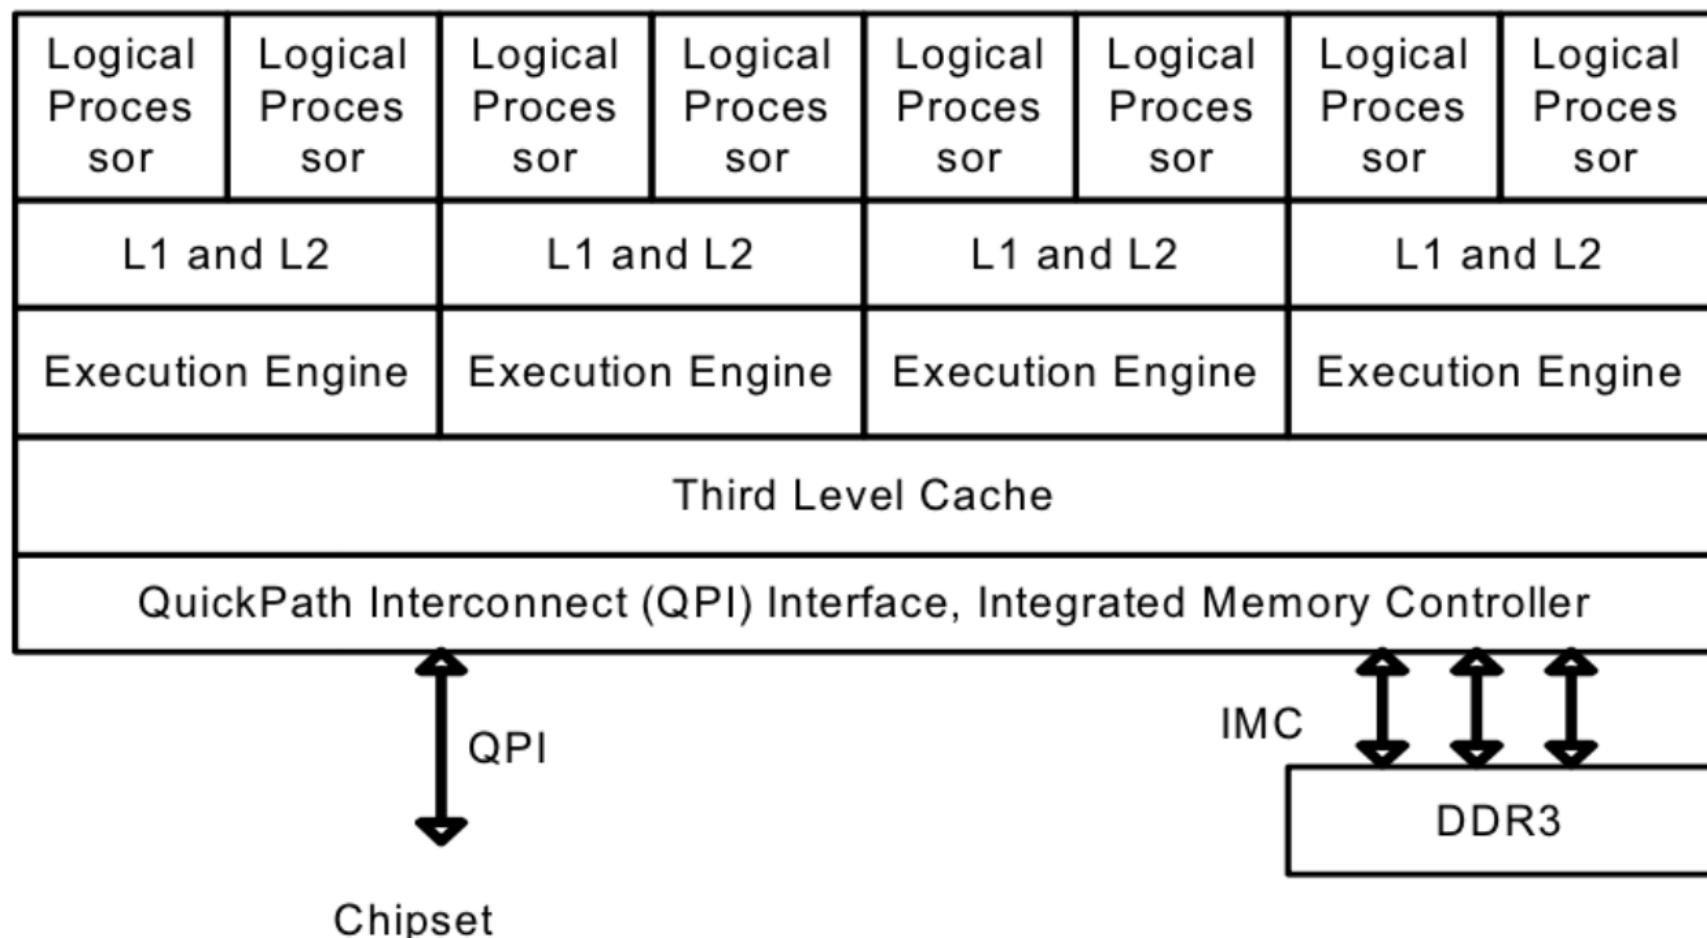
<http://www.slcentral.com>

Vector Units (SIMD)



Parallel Units on Chip

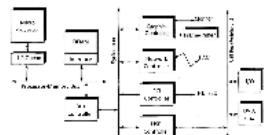
Intel Core i7 Processor



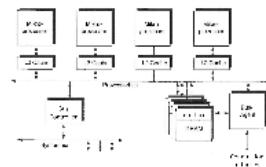
System Architectures

Replicating functional units

Workstation Architecture



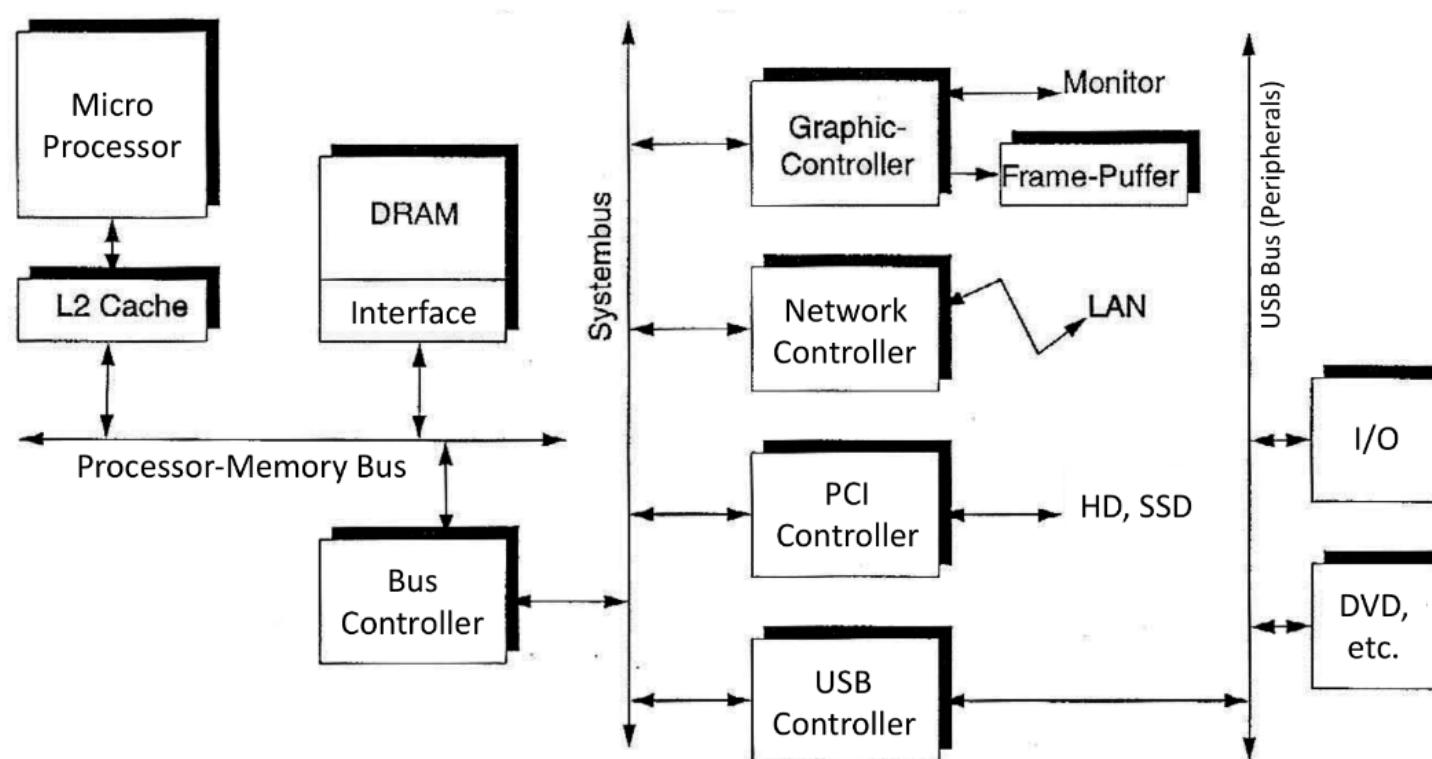
Multiprocessor Server Architecture



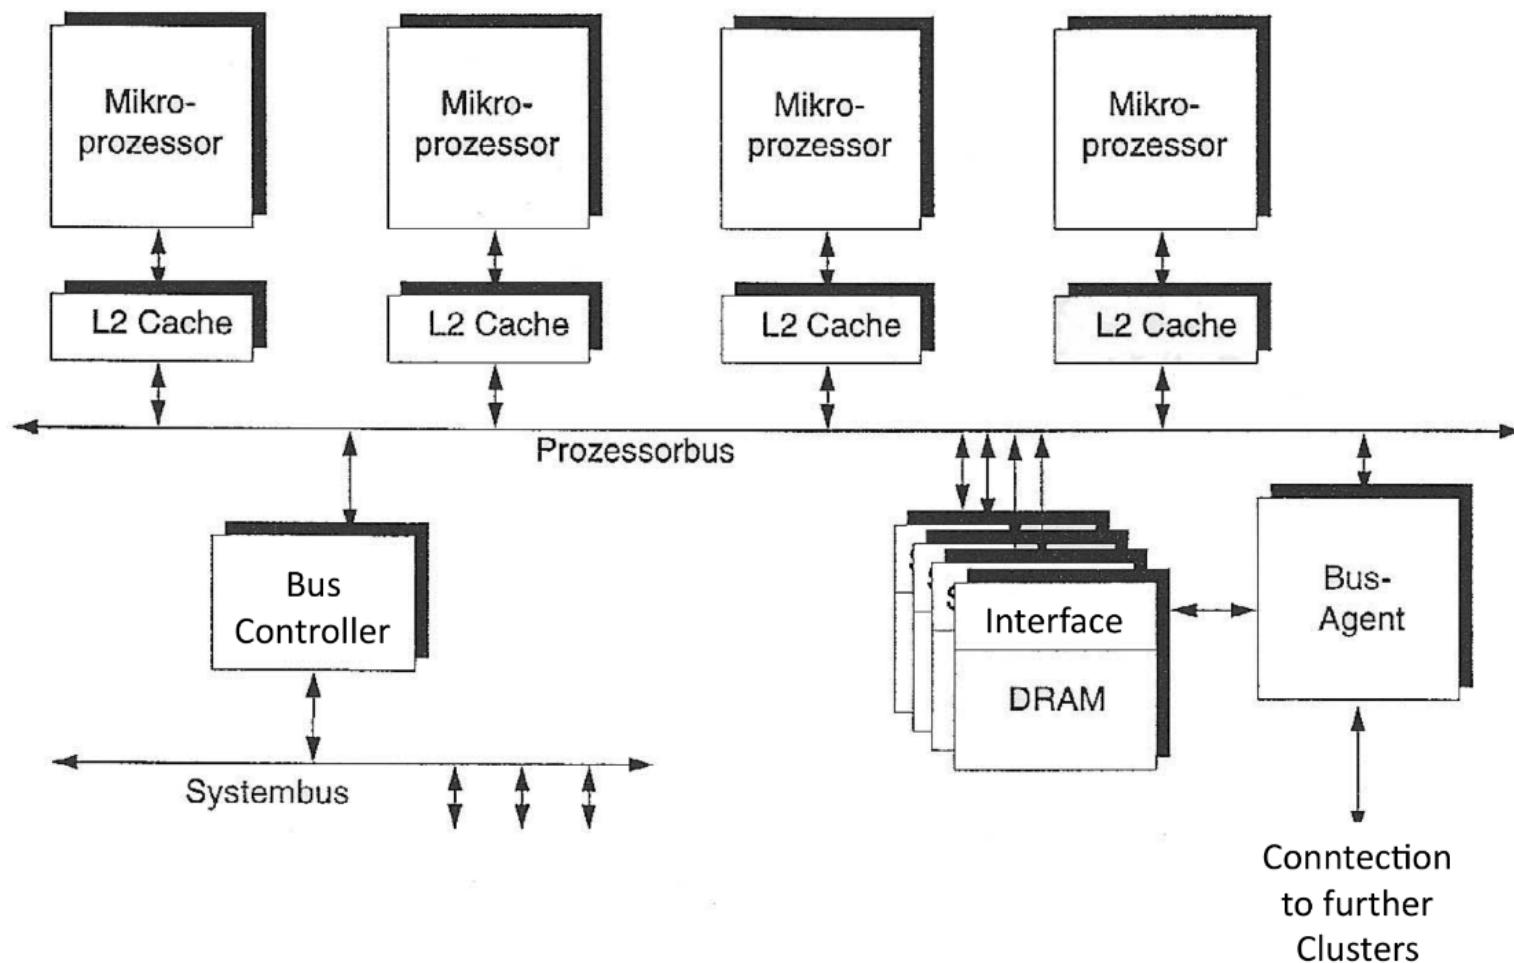
Latencies and Performance Indicators

Latency (ns) for various functional units						
	ALU	Branch	Cache	Divide	MUL	Memory
Processor	10	10	10	10	10	10
Memory	100	100	100	100	100	100
Cache	10	10	10	10	10	10
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	10	10
MUL	100	100	100	100	100	100
Divide	100	100	100	100	100	100
Branch	10	10	10	10	10	10
Cache	10	10	10	10	10	10
Memory	100	100	100	100	100	100
ALU	10	10	10	10	1	

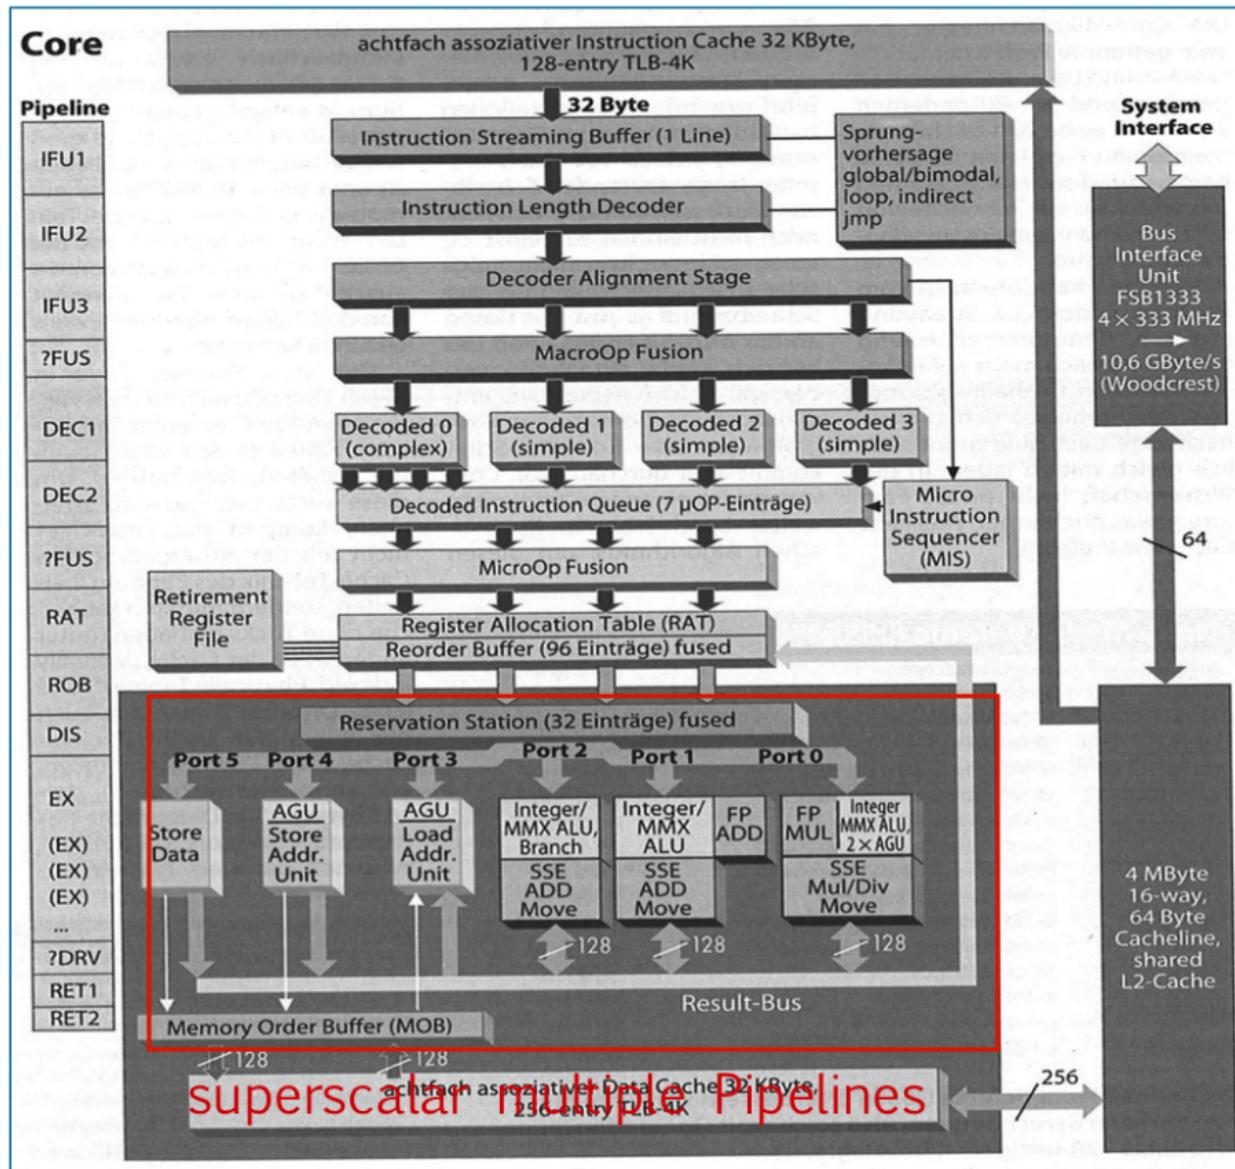
Worksstation Architecture



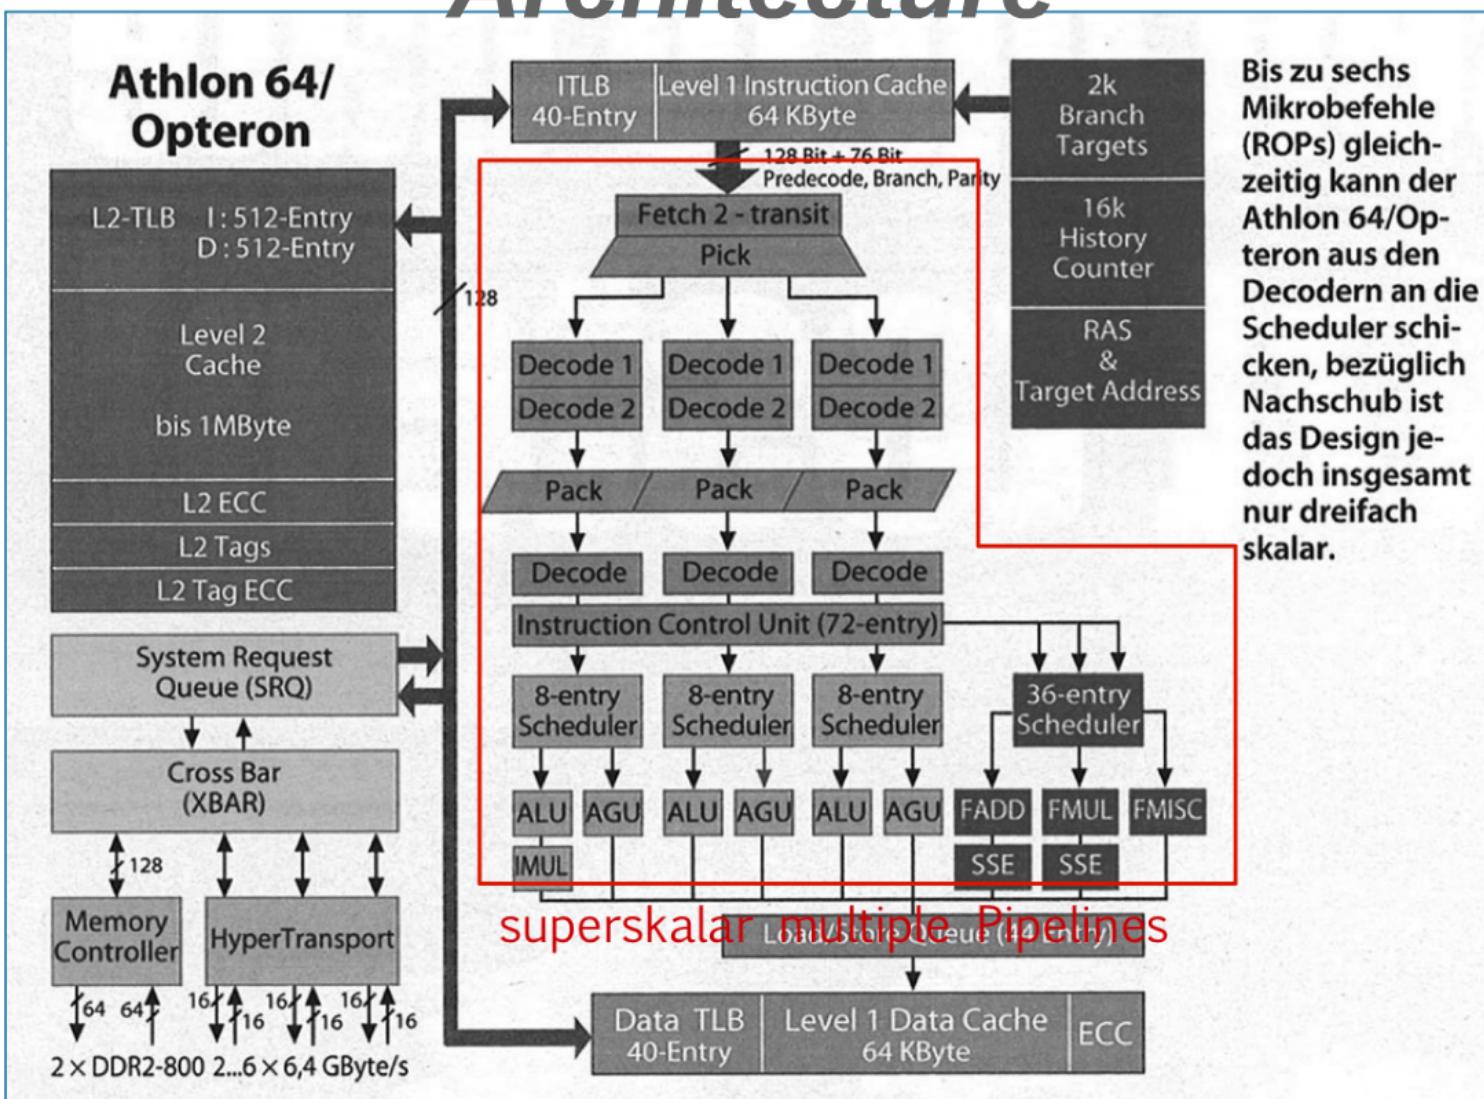
Multiprocessor Server Architecture



Intel Core Architecture



AMD Athlon/Opteron Architecture



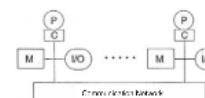
Bis zu sechs Mikrobefehle (ROPs) gleichzeitig kann der Athlon 64/Opteron aus den Decodern an die Scheduler schicken, bezüglich Nachschub ist das Design jedoch insgesamt nur dreifach skalar.

Multi-processor Architectures

A classification according to Flynn (1972)

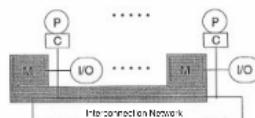


MIMD: Distributed Memory Multi-Computer



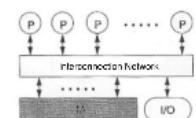
- Memory and periphery distributed
- No remote memory access
- Example: "Beowulf cluster"

MIMD: Virtual Shared Memory Multi-Processor



- Distributed memory accessible by all procs.
- Non-Uniform memory access (local access faster)
- CC-NUMA: Cache Coherent Non-Uniform Memory Access

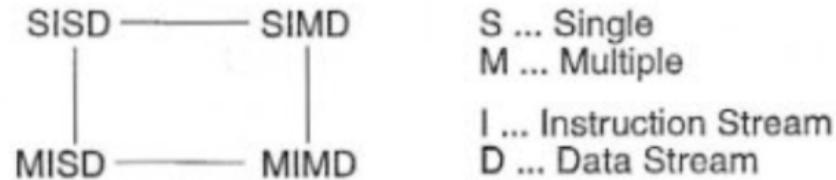
MIMD: Shared Memory Multi-Processor



- Memory and periphery accessible by all procs.
- Uniform Memory Access (UMA)
- Network: Bus or Crossbar

A classification according to Flynn (1972)

Very coarse, qualitative categorization of multiplicity of instruction and data streams



Examples:

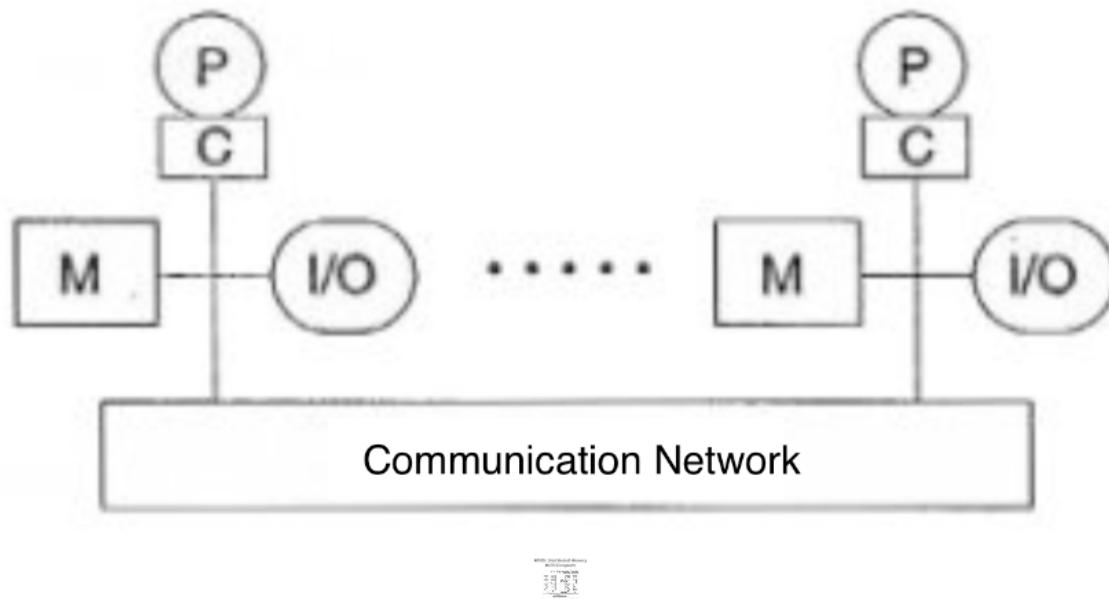
SISD: Monoprocessors, von-Neumann Architecture

MISD: ? (multithreading)

SIMD: Vector processors

MIMD: multi-processors, clusters, etc.

MIMD: Distributed Memory Multi-Computer

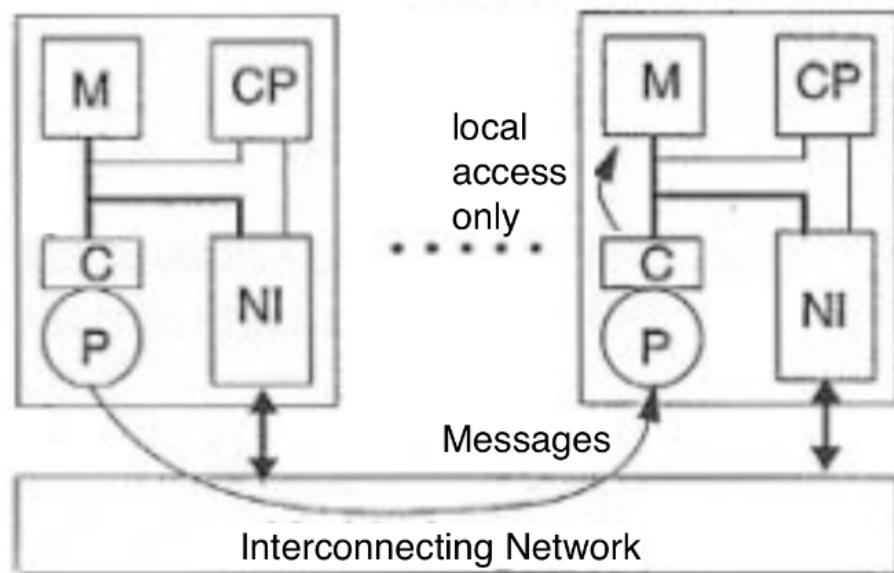


- Memory and periphery distributed
- No remote memory access
- Example: "Beowulf cluster"

MIMD: Distributed Memory Multi-Computer

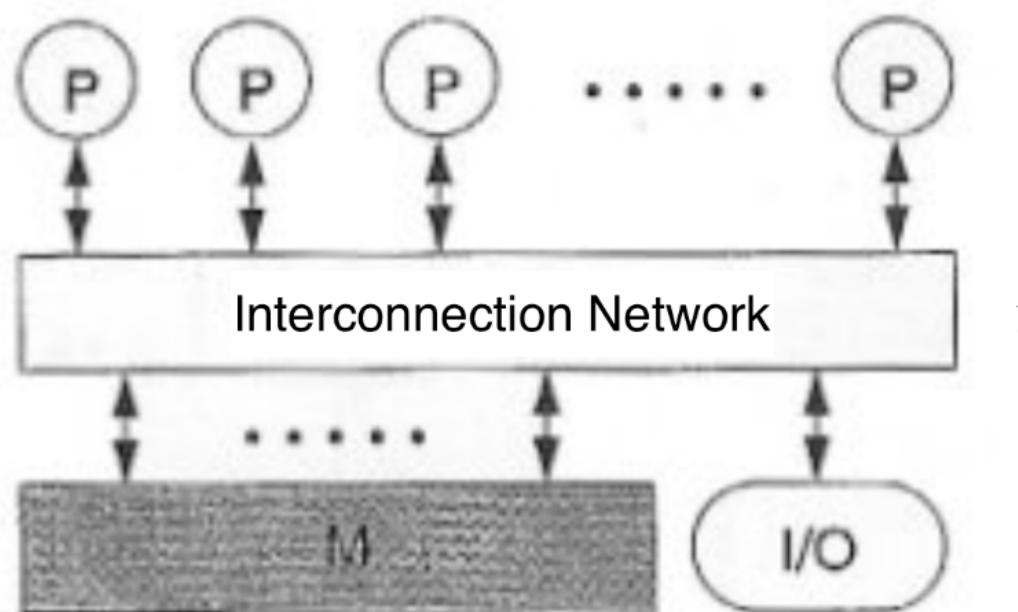
CP ... *Communication Processor*

NI ... *Network Interface*



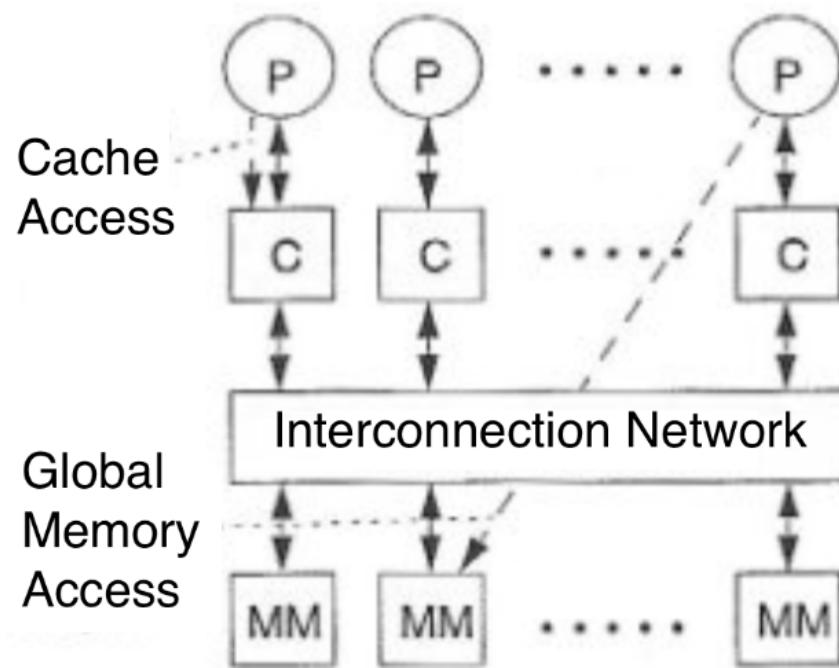
- Higher Bandwidth
- Low latency to local memory
- Dedicated message passing
- Custom Design

MIMD: Shared Memory Multi-Processor



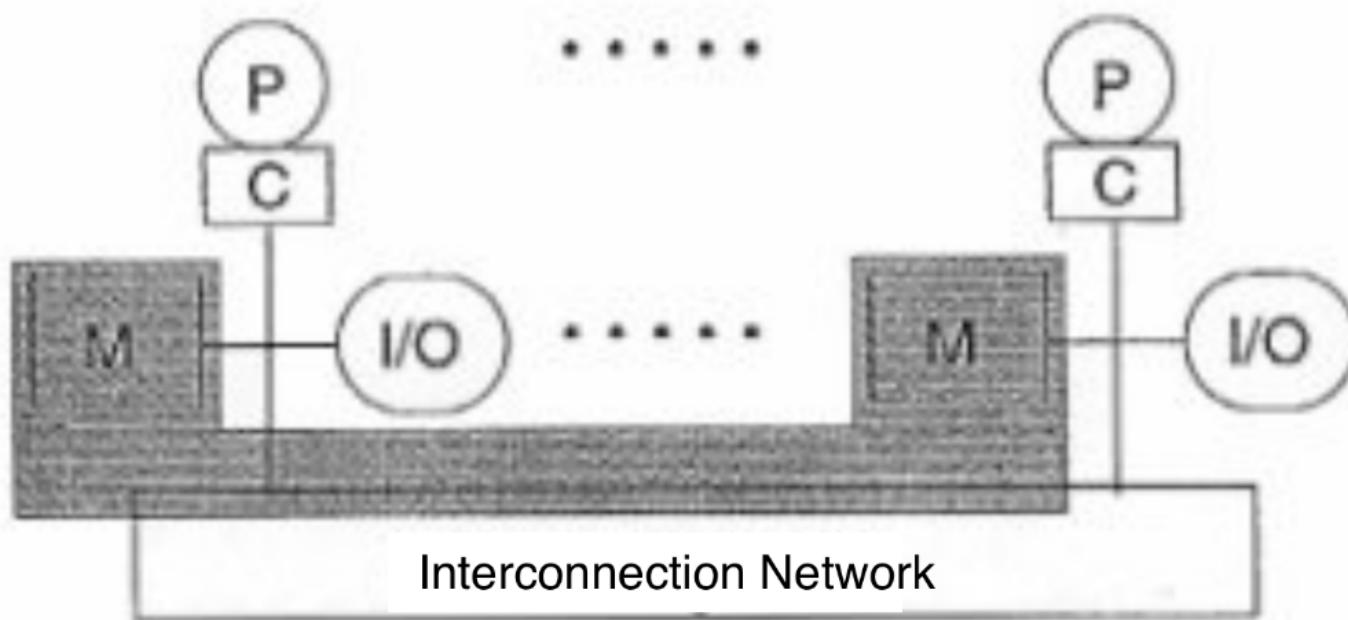
- Memory and periphery accessible by all procs.
- Uniform Memory Access (UMA)
- Network: Bus or Crossbar

MIMD: NUMA Shared Memory Multi-Processor



- Memory hierarchy (caches)
- Non-Uniform Memory Access (NUMA)
- Less Network traffic, faster (local) memory access
- Cache coherency required!

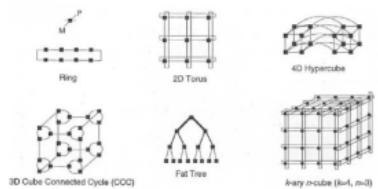
MIMD: Virtual Shared Memory Multi-Processor



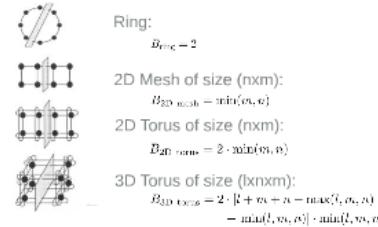
- Distributed memory accessible by all procs.
- Non-Uniform memory access (local access faster)
- CC-NUMA: Cache Coherent Non-Uniform Memory Access

Network Architectures

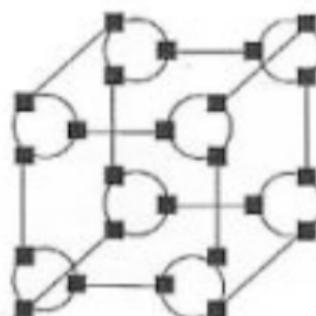
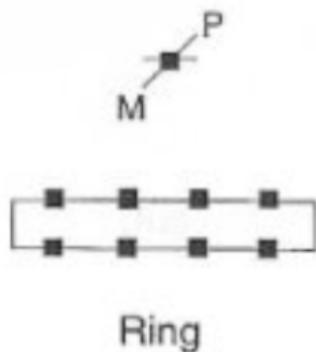
Topologies of static networks



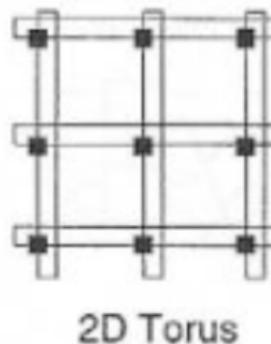
Determining the performance of a network: bisection bandwidth



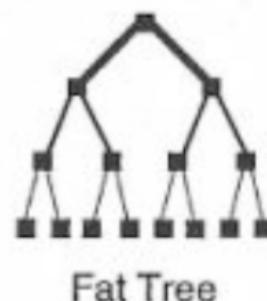
Topologies of static networks



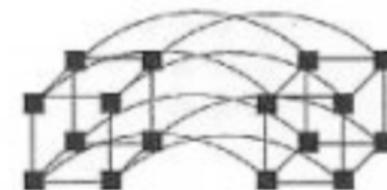
3D Cube Connected Cycle (CCC)



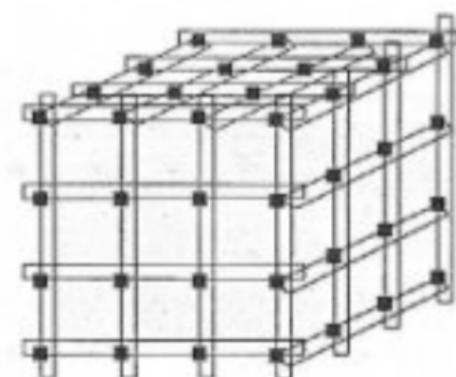
2D Torus



Fat Tree



4D Hypercube



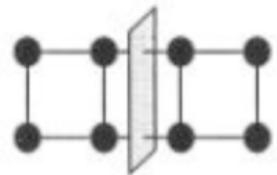
k -ary n -cube ($k=4$, $n=3$)

Determining the performance of a network: bisection bandwidth



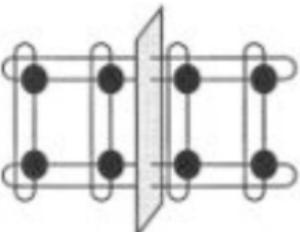
Ring:

$$B_{\text{ring}} = 2$$



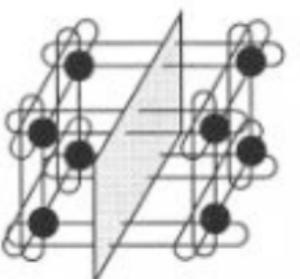
2D Mesh of size ($n \times m$):

$$B_{\text{2D mesh}} = \min(m, n)$$



2D Torus of size ($n \times m$):

$$B_{\text{2D torus}} = 2 \cdot \min(m, n)$$



3D Torus of size ($l \times n \times m$):

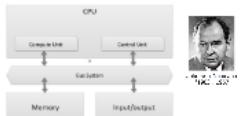
$$\begin{aligned} B_{\text{3D torus}} = & 2 \cdot [l + m + n - \max(l, m, n) \\ & - \min(l, m, n)] \cdot \min(l, m, n) \end{aligned}$$

Network Topology Characteristics

Topology	# Nodes	Degree (# ports)	min. path	bisection width
Ring	N	2	$\frac{N}{2}$	2
2D Torus	$N = k^2$	4	$2\frac{k}{2}$	$2k$
Hypercube (d-dim.)	$N = 2^d$	d	d	$\frac{N}{2}$
Tree	N	3	$2(\log_2 N - 1)$	1
k -ary n -cube	$N = k^n$	$2n$	$n\frac{k}{2}$	$2k^{n-1}$

Basics

Von Neumann Architecture



John von Neumann
1903 - 1957

Introduction

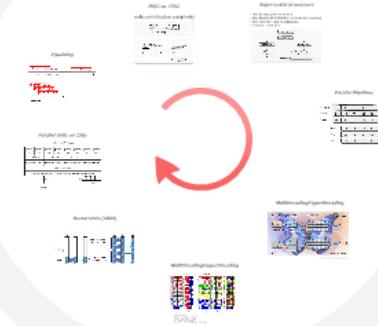
Why deal with Computer Architectures?

- Most programs achieve 10-20% of peak performance, why?
- How does a computer work, anyways?
- What makes the performance?



Acceleration Features

Parallelism on the processor level



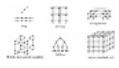
Computer Architectures

Scientific Computing
Jörn Behrens



Network Architectures

Typeologies of static networks



Determining the performance of a nation's backbone bandwidth

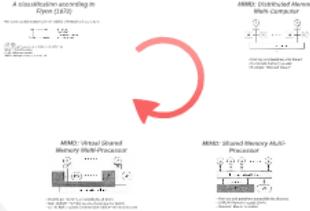


Multi-processor Architectures

A classification according to Open (2010)

Reproduced with permission of the copyright owner

Open, J., 2010. Multi-processor Architectures. In: J. Behrens, ed. Scientific Computing. Berlin, Heidelberg: Springer Berlin Heidelberg.



System Architectures

Replicating functional units

